

5ntiyayph

May 5, 2025

# 1 Boston Housing Price Prediction with Keras

## 1.0.1 Step 1: Import Required Libraries

```
[1]: # Data processing and visualization
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

# Sklearn tools
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score

# TensorFlow Keras
import tensorflow as tf
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Input
```

## 1.0.2 Step 2: Load and View the Dataset

```
[2]: df = pd.read_csv("datasets/boston_housing.csv")
df.head()
```

```
[2]:
```

	crim	zn	indus	chas	nox	rm	age	dis	rad	tax	ptratio	\
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	

	b	lstat	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4

```
4 396.90 5.33 36.2
```

### 1.0.3 Step 3: Split Features and Target

```
[3]: X = df.drop(['MEDV'], axis=1)
     y = df['MEDV']
```

### 1.0.4 Step 4: Scale the Features

```
[4]: scaler = StandardScaler()
     X_scaled = scaler.fit_transform(X)
```

### 1.0.5 Step 5: Split into Training and Testing Sets

```
[5]: X_train, X_test, y_train, y_test = train_test_split(
     X_scaled, y, test_size=0.2, random_state=42
)
```

### 1.0.6 Step 6: Build the Neural Network Model

```
[6]: model = Sequential([
     Input(shape=(X_train.shape[1],)),
     Dense(64, activation='relu'),
     Dense(32, activation='relu'),
     Dense(1)
])
```

### 1.0.7 Step 7: Compile the Model

```
[7]: model.compile(
     optimizer='adam',
     loss='mse',
     metrics=['mae']
)
```

### 1.0.8 Step 8: Train the Model

```
[8]: history = model.fit(
     X_train, y_train,
     epochs=50,
     validation_split=0.1,
     verbose=1
)
```

Epoch 1/50  
12/12 1s 14ms/step -  
loss: 581.5626 - mae: 22.2346 - val\_loss: 552.5659 - val\_mae: 21.9875

Epoch 2/50  
12/12 0s 3ms/step - loss:  
573.7921 - mae: 22.1359 - val\_loss: 512.7275 - val\_mae: 21.1069

Epoch 3/50  
12/12 0s 4ms/step - loss:  
500.3716 - mae: 20.4333 - val\_loss: 465.1572 - val\_mae: 20.0022

Epoch 4/50  
12/12 0s 3ms/step - loss:  
481.3070 - mae: 19.9221 - val\_loss: 403.9552 - val\_mae: 18.5018

Epoch 5/50  
12/12 0s 3ms/step - loss:  
411.9416 - mae: 18.0575 - val\_loss: 329.0402 - val\_mae: 16.5079

Epoch 6/50  
12/12 0s 4ms/step - loss:  
305.0258 - mae: 14.9884 - val\_loss: 245.1579 - val\_mae: 13.9974

Epoch 7/50  
12/12 0s 4ms/step - loss:  
213.0809 - mae: 12.1972 - val\_loss: 166.3586 - val\_mae: 11.0665

Epoch 8/50  
12/12 0s 3ms/step - loss:  
131.4219 - mae: 9.2341 - val\_loss: 107.1288 - val\_mae: 8.2478

Epoch 9/50  
12/12 0s 4ms/step - loss:  
104.0689 - mae: 8.0161 - val\_loss: 73.5620 - val\_mae: 6.1589

Epoch 10/50  
12/12 0s 4ms/step - loss:  
74.0127 - mae: 6.9324 - val\_loss: 57.3545 - val\_mae: 5.1876

Epoch 11/50  
12/12 0s 3ms/step - loss:  
60.5930 - mae: 6.2900 - val\_loss: 48.4163 - val\_mae: 4.6282

Epoch 12/50  
12/12 0s 3ms/step - loss:  
43.9501 - mae: 5.3162 - val\_loss: 43.9488 - val\_mae: 4.2788

Epoch 13/50  
12/12 0s 3ms/step - loss:  
32.3117 - mae: 4.5005 - val\_loss: 41.8687 - val\_mae: 4.1557

Epoch 14/50  
12/12 0s 5ms/step - loss:  
31.3455 - mae: 4.3628 - val\_loss: 40.8295 - val\_mae: 4.1188

Epoch 15/50  
12/12 0s 5ms/step - loss:  
27.0926 - mae: 3.9964 - val\_loss: 40.1871 - val\_mae: 4.1267

Epoch 16/50  
12/12 0s 4ms/step - loss:  
24.4301 - mae: 3.6270 - val\_loss: 39.1458 - val\_mae: 4.1493

Epoch 17/50  
12/12 0s 4ms/step - loss:  
23.1145 - mae: 3.7278 - val\_loss: 38.6148 - val\_mae: 4.1377  
Epoch 18/50  
12/12 0s 3ms/step - loss:  
25.4708 - mae: 3.7476 - val\_loss: 38.0797 - val\_mae: 4.1542  
Epoch 19/50  
12/12 0s 3ms/step - loss:  
25.0564 - mae: 3.6167 - val\_loss: 37.5187 - val\_mae: 4.1233  
Epoch 20/50  
12/12 0s 4ms/step - loss:  
21.3877 - mae: 3.4222 - val\_loss: 37.5324 - val\_mae: 4.0762  
Epoch 21/50  
12/12 0s 3ms/step - loss:  
22.5247 - mae: 3.4277 - val\_loss: 36.4104 - val\_mae: 3.9675  
Epoch 22/50  
12/12 0s 3ms/step - loss:  
19.5505 - mae: 3.2973 - val\_loss: 36.0621 - val\_mae: 3.9226  
Epoch 23/50  
12/12 0s 4ms/step - loss:  
20.2581 - mae: 3.4173 - val\_loss: 35.2850 - val\_mae: 3.8730  
Epoch 24/50  
12/12 0s 4ms/step - loss:  
21.1851 - mae: 3.3776 - val\_loss: 34.7377 - val\_mae: 3.8389  
Epoch 25/50  
12/12 0s 4ms/step - loss:  
18.5439 - mae: 3.1878 - val\_loss: 34.5024 - val\_mae: 3.8283  
Epoch 26/50  
12/12 0s 3ms/step - loss:  
21.4263 - mae: 3.1442 - val\_loss: 34.4287 - val\_mae: 3.7952  
Epoch 27/50  
12/12 0s 3ms/step - loss:  
17.6864 - mae: 3.0310 - val\_loss: 34.1468 - val\_mae: 3.7514  
Epoch 28/50  
12/12 0s 4ms/step - loss:  
15.6420 - mae: 2.9012 - val\_loss: 33.6757 - val\_mae: 3.7262  
Epoch 29/50  
12/12 0s 3ms/step - loss:  
17.0148 - mae: 3.0519 - val\_loss: 32.6068 - val\_mae: 3.6431  
Epoch 30/50  
12/12 0s 6ms/step - loss:  
15.1782 - mae: 2.8446 - val\_loss: 32.1881 - val\_mae: 3.5988  
Epoch 31/50  
12/12 0s 5ms/step - loss:  
17.2223 - mae: 3.0058 - val\_loss: 31.0931 - val\_mae: 3.5511  
Epoch 32/50  
12/12 0s 5ms/step - loss:  
14.0008 - mae: 2.8352 - val\_loss: 31.5389 - val\_mae: 3.5611

Epoch 33/50  
12/12 0s 9ms/step - loss:  
16.4224 - mae: 2.9407 - val\_loss: 30.9154 - val\_mae: 3.5132  
Epoch 34/50  
12/12 0s 5ms/step - loss:  
17.1202 - mae: 2.8992 - val\_loss: 30.8427 - val\_mae: 3.5197  
Epoch 35/50  
12/12 0s 5ms/step - loss:  
14.4388 - mae: 2.7627 - val\_loss: 30.0603 - val\_mae: 3.5058  
Epoch 36/50  
12/12 0s 5ms/step - loss:  
13.1486 - mae: 2.6630 - val\_loss: 29.8477 - val\_mae: 3.4819  
Epoch 37/50  
12/12 0s 4ms/step - loss:  
12.8587 - mae: 2.5735 - val\_loss: 29.6458 - val\_mae: 3.4632  
Epoch 38/50  
12/12 0s 5ms/step - loss:  
11.5893 - mae: 2.4936 - val\_loss: 28.6923 - val\_mae: 3.3993  
Epoch 39/50  
12/12 0s 4ms/step - loss:  
13.5457 - mae: 2.6462 - val\_loss: 27.9086 - val\_mae: 3.3470  
Epoch 40/50  
12/12 0s 4ms/step - loss:  
14.3147 - mae: 2.8455 - val\_loss: 28.4011 - val\_mae: 3.4203  
Epoch 41/50  
12/12 0s 5ms/step - loss:  
14.7176 - mae: 2.6634 - val\_loss: 27.8473 - val\_mae: 3.3859  
Epoch 42/50  
12/12 0s 5ms/step - loss:  
15.1011 - mae: 2.6999 - val\_loss: 27.5924 - val\_mae: 3.3982  
Epoch 43/50  
12/12 0s 5ms/step - loss:  
12.6673 - mae: 2.6357 - val\_loss: 27.6282 - val\_mae: 3.4089  
Epoch 44/50  
12/12 0s 4ms/step - loss:  
12.2337 - mae: 2.6344 - val\_loss: 26.2196 - val\_mae: 3.3371  
Epoch 45/50  
12/12 0s 5ms/step - loss:  
11.1612 - mae: 2.4315 - val\_loss: 26.5611 - val\_mae: 3.3573  
Epoch 46/50  
12/12 0s 4ms/step - loss:  
11.5243 - mae: 2.4337 - val\_loss: 26.4900 - val\_mae: 3.3527  
Epoch 47/50  
12/12 0s 4ms/step - loss:  
11.9757 - mae: 2.5080 - val\_loss: 25.4709 - val\_mae: 3.3489  
Epoch 48/50  
12/12 0s 4ms/step - loss:  
11.9704 - mae: 2.5175 - val\_loss: 25.3799 - val\_mae: 3.3569

```
Epoch 49/50
12/12          0s 5ms/step - loss:
11.5124 - mae: 2.4678 - val_loss: 25.4830 - val_mae: 3.3310
Epoch 50/50
12/12          0s 5ms/step - loss:
13.4999 - mae: 2.5512 - val_loss: 26.2756 - val_mae: 3.3284
```

### 1.0.9 Step 9: Predict on Test Set

```
[9]: y_pred = model.predict(X_test).flatten()
```

```
4/4          0s 11ms/step
```

### 1.0.10 Step 10: Evaluate Model Performance

```
[10]: mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)

print(f"MAE : {mae:.2f}")
print(f"MSE : {mse:.2f}")
print(f"R2 : {r2:.2f}")
```

```
MAE : 2.57
MSE : 16.16
R2 : 0.78
```

### 1.0.11 Step 11: Show Predictions vs Actual

```
[11]: print("Predicted Price vs Actual Price (First 5 Samples)")
for i in range(5):
    print(f"Predicted: {y_pred[i]:.2f} || Actual: {y_test.iloc[i]:.2f}")
```

```
Predicted Price vs Actual Price (First 5 Samples)
Predicted: 27.95 || Actual: 23.60
Predicted: 33.53 || Actual: 32.40
Predicted: 20.76 || Actual: 13.60
Predicted: 26.99 || Actual: 22.80
Predicted: 16.32 || Actual: 16.10
```

### 1.0.12 Step 12: Predict from User Input

```
[12]: def predict_from_user_input():
    print("-- Enter feature values to predict the house price --")

    feature_names = X.columns
    user_input = {}
```

```

for feature in feature_names:
    value = float(input(f"{feature}: "))
    user_input[feature] = value

user_df = pd.DataFrame([user_input])
user_scaled = scaler.transform(user_df)

pred = model.predict(user_scaled)
print(f"\nPredicted Price: ${pred[0][0]*1000:.2f}")

# Uncomment to use
# predict_from_user_input()


```

### 1.0.13 Step 13: Plot Actual vs Predicted Prices

```


[13]: plt.figure(figsize=(6, 4))
plt.scatter(y_test, y_pred, alpha=0.6)
plt.plot([min(y_test), max(y_test)], [min(y_test), max(y_test)], 'r--')
plt.xlabel("Actual Price")
plt.ylabel("Predicted Price")
plt.title("Actual vs Predicted Price")
plt.grid(True)
plt.show()

```



#### 1.0.14 Step 14: Plot Loss Curve


```
[14]: plt.plot(history.history['loss'], label='Train')
plt.plot(history.history['val_loss'], label='Validation')
plt.title("Loss Curve")
plt.xlabel("Epoch")
plt.ylabel("Loss")
plt.legend()
plt.grid(True)
plt.show()
```



#### 1.0.15 Step 15: Distribution of House Prices


```
[15]: sns.histplot(y_train, bins=20, kde=True, color='green')
plt.title("Distribution of House Prices (MEDV)")
plt.xlabel("Price")
plt.ylabel("Frequency")
plt.grid(True)
plt.show()
```





#### 1.0.16 Step 16: Correlation Heatmap of Features

```
[17]: plt.figure(figsize=(6,4))
sns.heatmap(df.corr(), annot=False, cmap='coolwarm')
plt.title("Feature Correlation Heatmap")
plt.show()
```



[ ]: