

# DL Code Explanations

Great! Let's now explain your **Boston Housing Price Prediction** practical in **detailed, viva-style language, step-by-step** — not just the code, but the *thinking behind every line*, so you can **confidently convince your examiner** you understand it fully.

---

## Main Goal

Predict the **price of houses in Boston** using **deep learning** (Neural Networks) with **Keras**.

- **Type of problem:** Regression (continuous value prediction — price, not class).
  - **Why neural network:** The relationship between features (like number of rooms, crime rate) and price is complex and possibly non-linear.
- 

## Step-by-Step Breakdown

---

### ♦ Step 1: Import Required Libraries

#### Why we do this?

We need tools to:

- Handle data → ``pandas`, `numpy``
- Plot graphs → ``matplotlib`, `seaborn``
- Preprocess data → ``StandardScaler``
- Split datasets → ``train_test_split``
- Evaluate performance → ``MAE`, `MSE`, `R²``
- Build models → ``tensorflow.keras``

✅ **What we're preparing for:** A full ML pipeline — from data to model training to evaluation.

---

### ♦ Step 2: Load and View the Dataset

python

```
df = pd.read_csv("datasets/boston_housing.csv")
df.head()
```

## Why?

We load the Boston Housing dataset which contains real data about houses (e.g., crime rate, number of rooms) and their prices.

- Each row = 1 house
- Each column = 1 feature (like `RM`, `CRIM`, `TAX`, etc.)
- Final column `MEDV` = price of the house (what we want to predict)

## ◆ Step 3: Split Features and Target

python

```
X = df.drop(['MEDV'], axis=1)
y = df['MEDV']
```

## Why?

We split the data into:

- `X` = features (input)
- `y` = target (output/label to predict)

 This is the basic requirement for **supervised learning**.

## ◆ Step 4: Feature Scaling

python

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

## Why scale features?

- Neural networks **converge faster** when input features are on a similar scale.
- If one feature (say, TAX) is in the 100s and another (say, CRIM) is <1, it causes **inefficient learning**.

- `StandardScaler`: makes mean = 0, std = 1.

## ♦ Step 5: Split into Train/Test Sets

python

```
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.2, random_state=42)
```

### Why split?

- **Training set:** Used to teach the model.
- **Testing set:** Used to **evaluate how well** the model generalizes to new data.
- `test_size=0.2` → 20% test, 80% train.
- `random_state` → ensures same split every time.

## Neural Network Modeling with Keras

## ♦ Step 6: Build the Model

python

```
model = Sequential([
    Input(shape=(X_train.shape[1],)),
    Dense(64, activation='relu'),
    Dense(32, activation='relu'),
    Dense(1)
])
```

### What's happening?

- **Sequential:** Linear stack of layers.
- **Input layer:** Number of inputs = number of features (13).
- **Hidden layers:**
  - 64 neurons (ReLU): Learns patterns.
  - 32 neurons (ReLU): Deeper representation.
- **Output layer:** 1 neuron (price), **no activation** → because this is **regression**.

 **ReLU:** Adds non-linearity so the model can learn complex relationships.

## ◆ Step 7: Compile the Model

```
python

model.compile(optimizer='adam', loss='mse', metrics=['mae'])
```

### ⚙️ Why these settings?

- **Loss = MSE** (Mean Squared Error): Perfect for regression; penalizes large errors.
- **Optimizer = Adam**: Efficient, adaptive learning.
- **Metric = MAE**: Easy to understand (e.g., model is off by ₹20,000 on average).

## ◆ Step 8: Train the Model

```
python

history = model.fit(
    X_train, y_train,
    epochs=50,
    validation_split=0.1,
    verbose=1
)
```

### 📚 Why?

- Model learns from training data in **50 passes (epochs)**.
- **10%** of training data used for validation — helps track overfitting.
- Training adjusts model weights to **minimize the loss (MSE)**.

## ◆ Step 9: Predict on Test Set

```
python

y_pred = model.predict(X_test).flatten()
```

### 🧠 Why?

- We use the trained model to make predictions on **unseen data**.
- `flatten()` converts output from shape `(n,1)` to `(n,)` for ease.

## ◆ Step 10: Evaluate the Model

```
python
```

```
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
```

## Why multiple metrics?

Metric	What it tells
MAE	Avg error (e.g., ₹3.5k)
MSE	Heavily penalizes big errors
R <sup>2</sup>	Variance explained by model (closer to 1 = better)

## ♦ Step 11: Show Predictions vs Actual

```
python
```

```
for i in range(5):
    print(f"Predicted: {y_pred[i]:.2f} || Actual: {y_test.iloc[i]:.2f}")
```

## Why?

Manually check if predictions make sense — good to show in the viva.

## ♦ Step 12: Predict from User Input

```
python
```

```
# function takes real inputs from user, scales, and predicts
```

## Why include this?

- Demonstrates that the model can be used in **real-time applications** (e.g., house price tool on a website).
- Shows understanding of **deployment-readiness**.

## ♦ Step 13: Plot Actual vs Predicted Prices

## Why this graph?

- X-axis: Actual price
- Y-axis: Predicted price

- Red diagonal = perfect predictions
- Points close to line = accurate model

👁️ Helps **visually judge** performance.

---

## ♦ Step 14: Plot Loss Curve

 Why this?

- Shows how **loss decreases over epochs**.
  - Helps check:
    - Is model learning? (loss should drop)
    - Overfitting? (validation loss much higher than training loss)
- 

## ♦ Step 15: Plot Distribution of House Prices

 Why this?

- Shows how the target (`MEDV`) is spread.
  - Helps understand if data is skewed or balanced.
  - Useful before choosing model/loss function.
- 

## ♦ Step 16: Plot Correlation Heatmap

 Why this?

- See how each feature correlates with others and with `MEDV`.
  - Helps in:
    - Feature selection
    - Understanding relationships
- 

## Quick Summary Table for Viva

Step	Goal	Tools
Data Prep	Clean, split, scale	pandas, sklearn
Modeling	Build & train NN	Keras Sequential

Step	Goal	Tools
Evaluation	Predict, score	MAE, MSE, R <sup>2</sup>
Visualization	Understand performance	matplotlib, seaborn

Do you want me to now prepare a **last-minute revision version** of this explanation — just 1–2 lines per step?

Here's a full explanation of the **IMDB Sentiment Classification Using Keras**, explained like in the ``boston.pdf`` file — focusing on **what is happening, why it's done, and how it works**, not on the code.

## **IMDB Sentiment Classification – What and Why?**

The goal is to **predict whether a movie review is positive or negative**, based on the text of the review. This is a **binary classification** problem. We're using **deep learning (Keras)** to build a model that understands patterns in text and predicts sentiment.

### ◆ **Step-by-Step Explanation**

#### ◆ **Step 1: Import Libraries**

- We load libraries to:
  - Handle data (``pandas``)
  - Preprocess text (``Tokenizer``, ``pad_sequences``)
  - Build deep learning models (``Sequential``, ``Embedding``, etc.)
  - Evaluate results (``sklearn``, ``seaborn``)

✅ **Why?** These are all essential tools for data processing and model building.

#### ◆ **Step 2: Load Dataset**

- We load ``imdb_master.csv`` which contains movie reviews.
- We specify a special encoding (``ISO-8859-1``) to handle special characters properly.

✅ **Why?** So the reviews are readable and correctly interpreted.

#### ◆ **Step 3: Drop Unnecessary Columns**

- Remove irrelevant columns like file names or indices that don't help with sentiment.

✅ **Why?** To clean the data and keep only what we need.

---

#### ♦ **Step 4: Explore Label Distribution**

- Count how many reviews are labeled positive, negative, or unsupervised.

✅ **Why?** To understand the data balance and remove noise.

---

#### ♦ **Step 5: Filter Dataset**

- Keep only **train** type reviews with **pos** or **neg** labels.
- Ignore unsupervised reviews because they don't help train the model.

✅ **Why?** We need labeled data to train a supervised model.

---

#### ♦ **Step 6: Separate Features and Labels**

- ``texts`` → the actual reviews
- ``y`` → the sentiments (pos or neg)

✅ **Why?** We need to train on inputs (``texts``) and outputs (``y``).

---

#### ♦ **Step 7: Encode Labels**

- Convert "pos" → 1 and "neg" → 0 using ``LabelEncoder``.

✅ **Why?** ML models don't understand words like "pos" or "neg" — only numbers.

---

#### ♦ **Step 8: Tokenize the Text**

- Create a vocabulary of the 10,000 most frequent words.
- Turn text into sequences of integers (each word becomes a number).

✅ **Why?** Neural networks don't understand text — we convert it to numbers.

---

#### ♦ **Step 9: Convert Text to Sequences**



- Every review becomes a list of word indexes.

✅ **Why?** These sequences will be input to the model.

---

## ♦ Step 10: Pad Sequences

- Make all reviews the same length (200 words). Shorter ones are padded with 0s.

✅ **Why?** Neural networks require fixed-size inputs.

---

## ♦ Step 11: Train-Test Split

- Split data into:
  - **Training set** (80%) to learn patterns.
  - **Testing set** (20%) to check performance.

✅ **Why?** So we know how well the model works on unseen data.

---

## ♦ Step 12: Build the Model

- Create a neural network with these layers:
  - **Embedding Layer:** Turns word indices into dense vectors.
  - **GlobalAveragePooling1D:** Averages word vectors to get a single sentence vector.
  - **Dense Layer + Dropout:** Learns useful patterns while avoiding overfitting.
  - **Output Layer (sigmoid):** Gives a probability of the review being positive.

✅ **Why?** The model learns to convert text into sentiment prediction.

---

## ♦ Step 13: Compile the Model

- Set:
  - **Loss function:** ``binary_crossentropy`` for binary classification.
  - **Optimizer:** ``adam`` for fast and stable training.
  - **Metric:** ``accuracy`` to monitor how often it's right.

✅ **Why?** This prepares the model for training.

---

### ♦ Step 14: Train the Model

- Run the training for 5 rounds (**epochs**) with 64 reviews at a time (**batch size**).
- Also check performance on the test set during training.

✅ **Why?** The model gradually improves by adjusting weights to minimize error.

---

### ♦ Step 15: Evaluate Model Accuracy

- Measure how well the model performs on the test set.

✅ **Why?** To check how accurate the predictions are.

---

### ♦ Step 16: Make Predictions

- Use the trained model to predict sentiment of the first 5 test reviews.

✅ **Why?** To test if the model gives good output on real examples.

---

### ♦ Step 17: Display Predictions vs Actual Labels

- Print whether the prediction was correct for each review.

✅ **Why?** To manually verify performance.

---

### ♦ Step 18: Visualize Training Label Distribution

- Plot how many positive vs negative reviews were in training data.

✅ **Why?** To check for any imbalance that could affect training.

---

## Bonus Visualizations (Not Numbered)

### ♦ Confidence Bar Plot

- Shows how confident the model is for 10 predictions using probability scores.

✅ **Why?** High confidence = strong prediction, Low confidence = uncertain.

## ◆ Confusion Matrix

- Shows how many:
  - Positives were predicted correctly
  - Negatives were predicted correctly
  - Mistakes made

✅ **Why?** Helps understand what types of errors the model makes.

## ◆ ROC Curve

- Plots True Positive Rate vs False Positive Rate.
- The curve shows how well the model distinguishes between classes.
- AUC (Area Under Curve) closer to 1 means better performance.

✅ **Why?** To visualize model quality beyond just accuracy.

## ✅ Summary

You trained a neural network to understand text reviews and predict if they are positive or negative by:

1. Cleaning and converting text into numeric form
2. Building a lightweight but powerful neural model
3. Training it on labeled data
4. Evaluating its performance using predictions, confusion matrix, and ROC curve

Would you like a **side-by-side visual comparison** of Boston vs IMDB process or a **printable diagram**?

Thanks for sharing the file. Based on the document and your request, here's a **very detailed explanation** of what's happening in the assignment, focusing on **what is happening, why it's done, and how it works** — all without diving deep into the code. The model used is a **Convolutional Neural Network (CNN)** applied to the **Fashion MNIST** dataset, a classic deep learning task.

## What Is Happening: An Overview of the Assignment

This assignment builds, trains, and evaluates a **Convolutional Neural Network (CNN)** to recognize images of clothing items using the **Fashion MNIST** dataset. The key steps include:

1. **Loading the dataset:** The Fashion MNIST dataset consists of 70,000 grayscale images of 10 fashion categories (e.g., shirts, trousers, shoes).
  2. **Data preprocessing:** Normalization and reshaping to make the data suitable for input into a CNN.
  3. **Building the CNN architecture:** Stacking layers that learn increasingly complex features from raw image data.
  4. **Compiling and training the model:** Teaching the CNN to minimize the error in classifying images correctly.
  5. **Evaluating the model:** Testing the trained CNN on unseen images and analyzing its performance.
  6. **Visualizing predictions:** Showing how well the model performs by displaying predicted labels for sample test images.
- 

## ? Why Each Step Is Done: Purpose Behind Each Stage

### 1. Loading Fashion MNIST

- **Why?** The Fashion MNIST dataset is a drop-in replacement for the older MNIST digits dataset and presents a more challenging classification problem.
  - It helps train models that need to **distinguish between similar-looking clothing items**, like shirts vs. tops.
- 

### 2. Preprocessing the Data

- **Normalization:** The pixel values (originally from 0 to 255) are scaled down to 0 to 1.
    - **Why?** Helps the neural network learn faster and more reliably by bringing data to a common scale.
  - **Reshaping:** Images are reshaped to add a **channel dimension** ( $28 \times 28 \rightarrow 28 \times 28 \times 1$ ).
    - **Why?** CNNs expect 3D input (height, width, channels). Grayscale images have one channel.
- 

### 3. Creating the CNN Architecture

CNNs are inspired by the human visual system. Here's a breakdown:

### ► Convolutional Layers

- **What?** These layers slide filters (or kernels) over the image to detect patterns (edges, textures).
- **Why?** To learn low-level and high-level features automatically without manual feature extraction.
- **How?** Each filter highlights specific patterns, creating **feature maps** that get passed to the next layer.

### ► Activation Function (ReLU)

- **What?** Applies a non-linearity ( $\max(0, x)$ ) to introduce non-linear properties.
- **Why?** Without non-linearity, the CNN would just be a linear function — not enough to learn complex patterns.
- **How?** ReLU keeps only positive values, helping the network focus on activated neurons.

### ► MaxPooling Layer

- **What?** Reduces the spatial size of the feature maps (downsampling).
- **Why?** Makes the model faster and reduces overfitting by retaining important features and discarding noise.
- **How?** Takes the maximum value in a small window (e.g., 2x2) over the feature map.

### ► Flatten Layer

- **What?** Converts the 2D feature maps into a 1D vector.
- **Why?** Fully connected layers (Dense layers) require a 1D input.
- **How?** Flattens all values into a long vector while maintaining order.

### ► Dense (Fully Connected) Layers

- **What?** Neurons are connected to every output from the previous layer.
- **Why?** These layers make the final prediction by combining learned features.
- **How?** Weighted sum of all inputs followed by an activation function (like ReLU or Softmax).

## ► Dropout Layer

- **What?** Randomly deactivates some neurons during training.
- **Why?** Prevents overfitting by ensuring the network doesn't rely too heavily on any one feature.
- **How?** Temporarily drops a % of neurons during training steps.

## ► Softmax Layer (Output Layer)

- **What?** Outputs probabilities for each of the 10 clothing classes.
  - **Why?** Needed for multi-class classification — model picks the class with the highest probability.
  - **How?** Converts raw output (logits) into probabilities that sum to 1.
- 

## 4. Compiling the Model

- **Loss Function:** ``sparse_categorical_crossentropy`` — used for multi-class classification with integer labels.
- **Optimizer:** ``adam`` — an efficient version of gradient descent.
- **Metrics:** Accuracy — shows how often predictions are correct.

### Why?

- Loss function tells the model how bad it's doing.
  - Optimizer updates the model to reduce the loss.
  - Metrics are just for our evaluation and tracking.
- 

## 5. Training the Model

- The CNN is shown training images repeatedly (epochs), and each time it adjusts its filters and weights to reduce the error.
  - **Why?** Training allows the model to "learn" by updating itself based on the mistakes it made during prediction.
  - **How?** Backpropagation: error is calculated → propagated back → weights updated → process repeats.
- 

## 6. Evaluating the Model

- The trained model is tested on images it has never seen before.
  - **Why?** To see how well it generalizes to new data.
  - **How?** Accuracy and loss are calculated on test data.
- 

## 7. Making and Visualizing Predictions

- A few test images are shown along with the model's predicted labels.
  - **Why?** Helps visually confirm that the model is working correctly and not just giving good numbers.
  - **How?** Predicted class = class with the highest probability from the Softmax output.
- 

## Extra Details to Keep in Mind

- **Epochs and Batch Size:**
    - An **epoch** is one full pass through the entire training data.
    - **Batch size** determines how many samples are processed before updating the model.
  - **CNNs are powerful** because:
    - They can capture spatial hierarchy (edges → shapes → objects).
    - They reduce the number of parameters using local connections (compared to fully connected networks).
    - They are translation invariant (they can detect patterns regardless of their location).
- 

## How the Model Learns (In Human Terms)

Imagine you're learning to recognize clothes:

- First, you look at a shirt and notice the **collar shape** or **button pattern** (like the convolutional filters).
- Next, you ignore parts like wrinkles (like max pooling filtering out noise).
- You make notes (memory), refine your guesses (like updating weights), and repeat the process with more examples.

Eventually, when someone shows you a new shirt, you immediately recognize it — that's exactly what CNNs are trained to do.

---

Would you like a **diagram** that explains how the CNN layers work visually on a Fashion MNIST image?

Is this conversation helpful so far?